

# Working with PyTorch Tensors

---



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Tensor operations using PyTorch tensors**

**Interoperability with NumPy**

**Understanding PyTorch support for GPUs and CUDA**

**Working with tensors on GPU-enabled devices**

# Tensors in PyTorch

---

# Tensor

The central unit of data in PyTorch. A tensor consists of a set of primitive values shaped into an array of any number of dimensions.

# Data Is Represented as Tensors



**Scalars are 0-D tensors**

**3, 6.7, “a”**

# Data Is Represented as Tensors



**Vectors** are **1-D** tensors

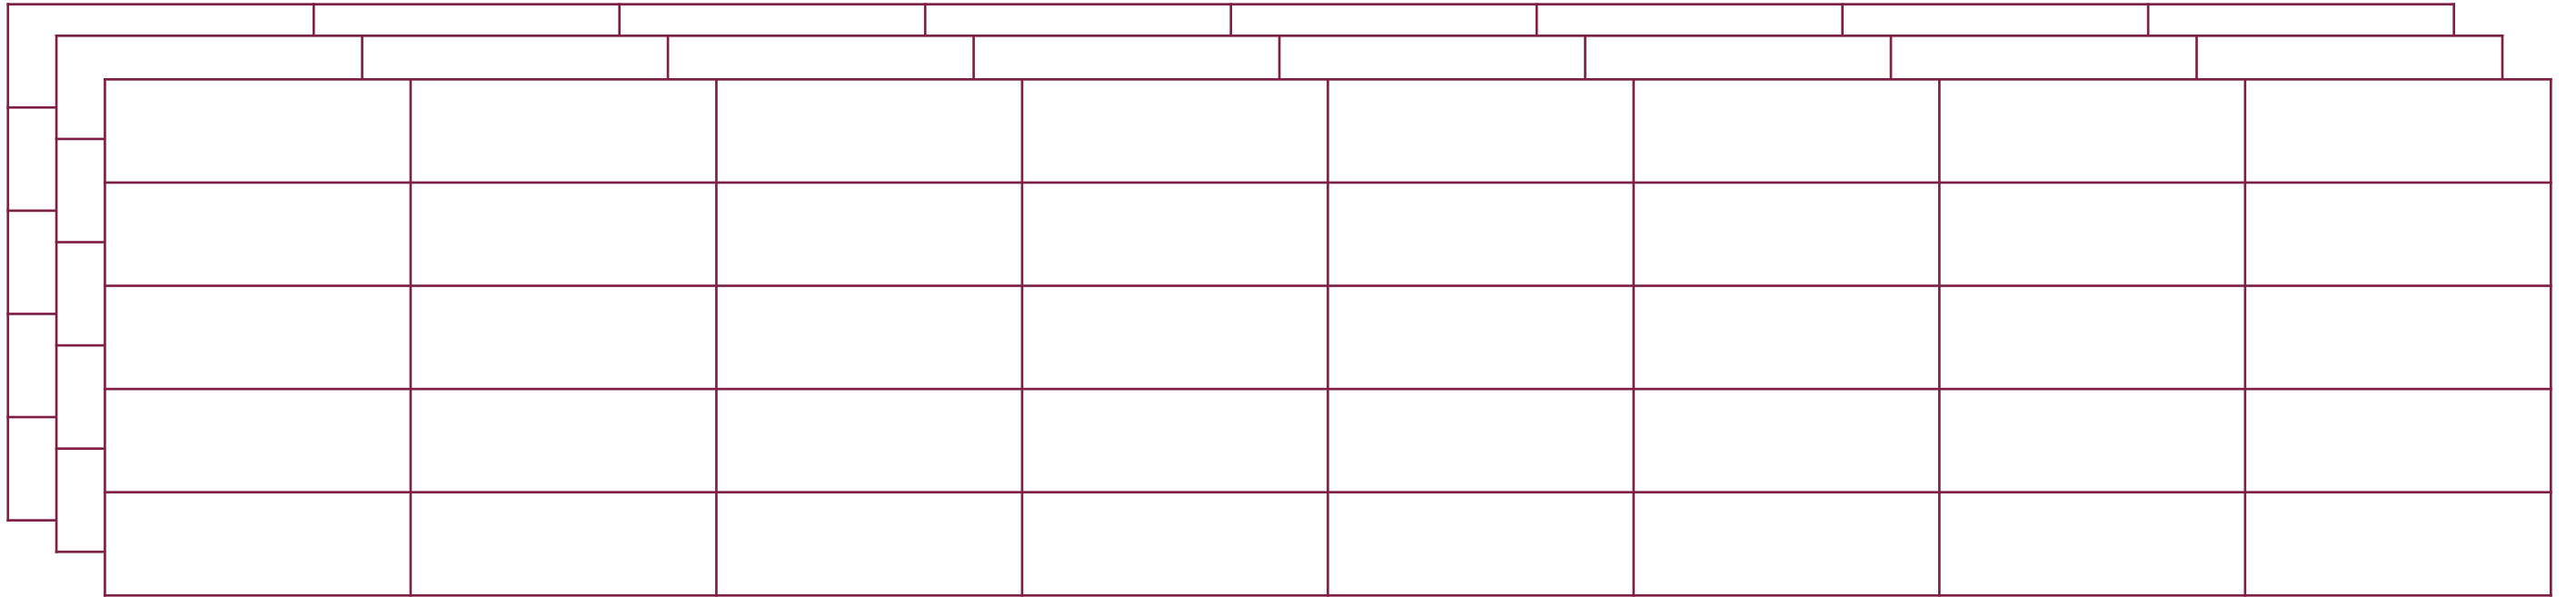
**[1, 3, 5, 7, 9]**

# Data Is Represented as Tensors


**Matrices** are **2-D** tensors

**[[1, 3, 5],  
[7, 9, 11]]**

# Data Is Represented as Tensors



**N-Dimensional matrices** are **N-D** tensors

**[[[1, 2], [3, 4], [5, 6],  
[7, 8], [9, 10], [11, 12]]]**



PyTorch Tensors have been architected  
to make optimal use of GPUs for  
massively parallel computations

Demo

**Operations using PyTorch tensors**

Demo

**Conversions between PyTorch and  
NumPy**

# PyTorch, CUDA, and GPUs

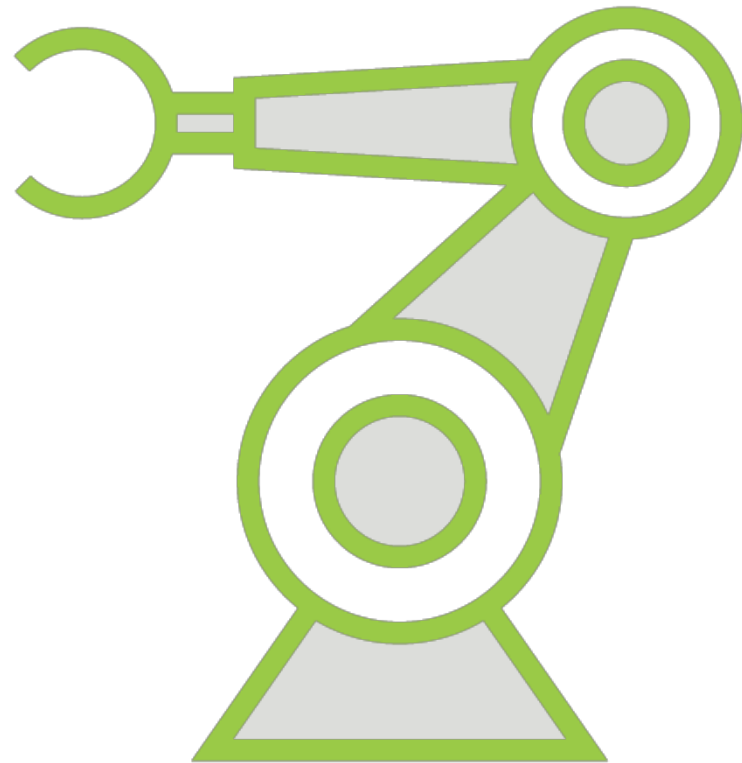
---

# GPU (Graphics Processing Unit)

Specialized chips with highly parallel architecture that makes them an order of magnitude faster than CPUs for some deep learning applications

PyTorch Tensors have been architected  
to make optimal use of GPUs for  
massively parallel computations

# GPUs for ML

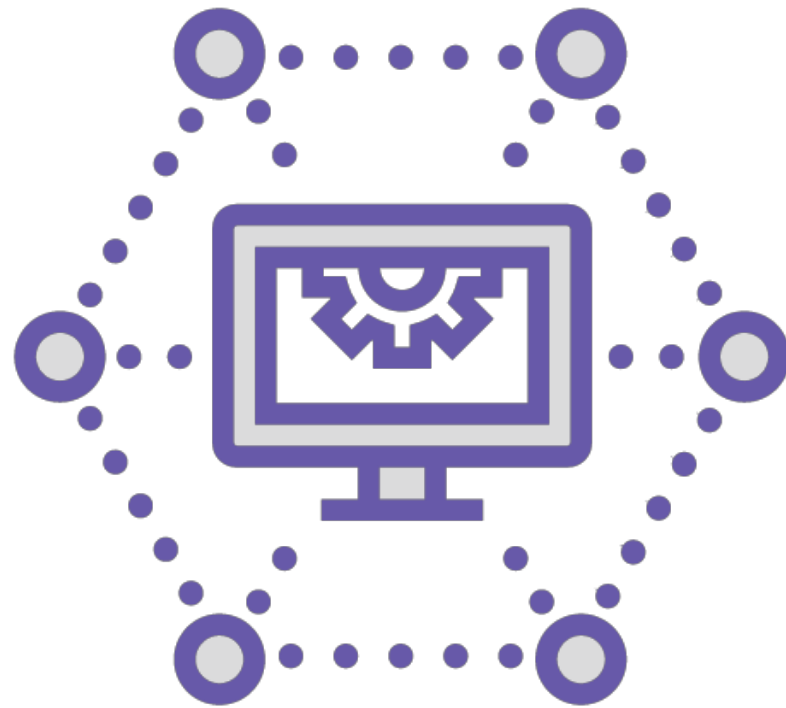


**Usage of GPUs has gone far beyond video/graphics processing**

**Widely used in Big Data and Machine Learning applications**

**Speedup of 10-50X where parallelization yields big wins**

# CUDA



**Nvidia is a major maker of GPUs**

**Devised CUDA, a parallel computing platform and API**

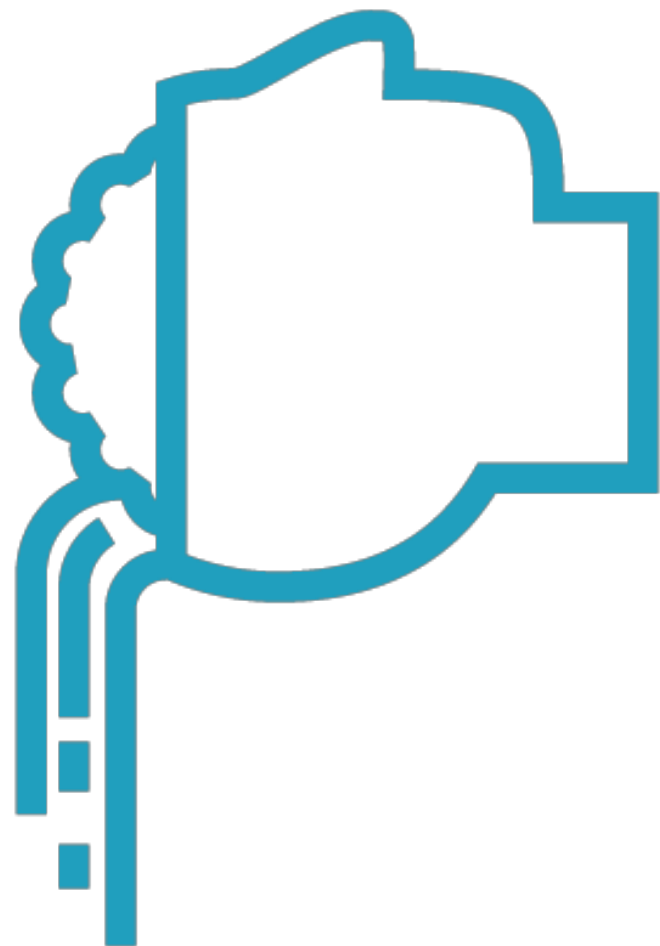
**A standard for general purpose (non-graphics) users of GPUs**

**Initially acronym for “Compute Unified Device Architecture”**

**Now a standalone term, not an acronym**



# CUDA and PyTorch

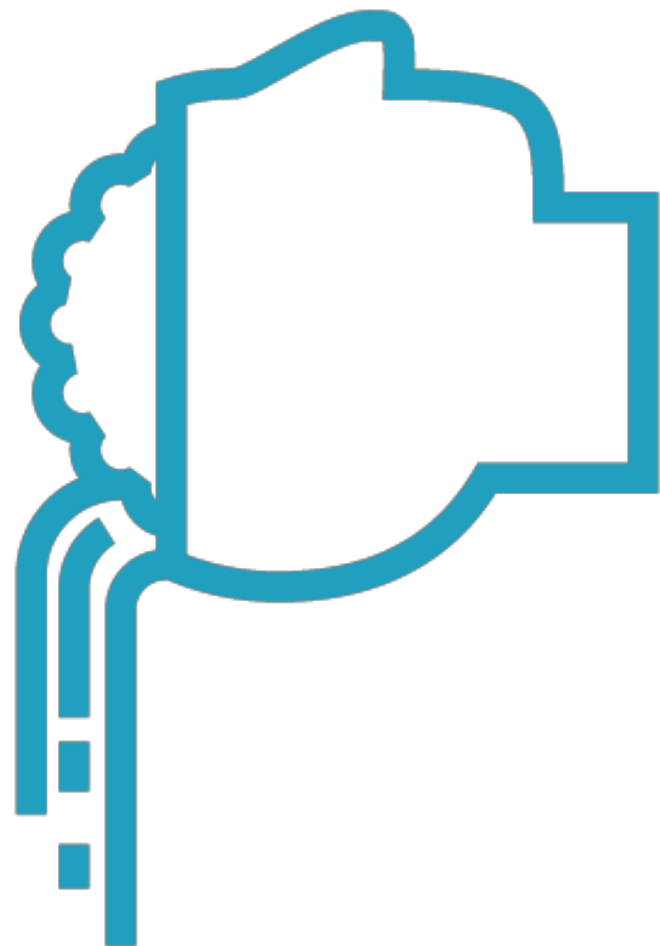


**Developers can write CUDA-compliant code**

**Code must be understood by CUDA-aware framework (e.g. PyTorch)**

**If CUDA-enabled GPUs are available, speedup will automatically occur**

# CUDA and PyTorch



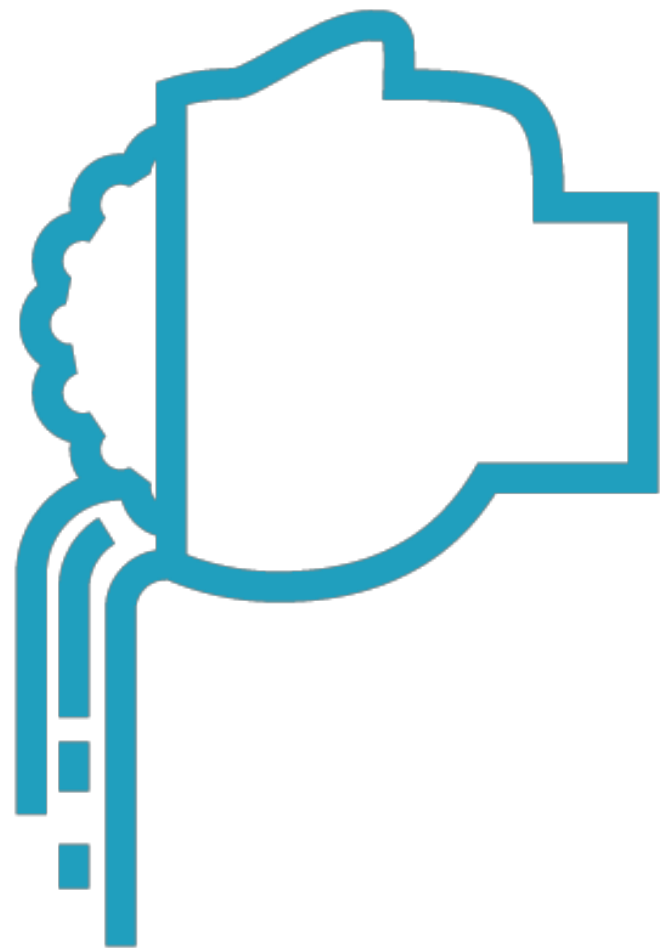
**`torch.cuda` for CUDA operations**

**Special tensor types for CUDA e.g.**  
**`torch.cuda.FloatTensor`**

**`torch.cuda.device` to select GPU**

**Tracks currently selected GPU and  
creates tensors on it**

# CUDA and PyTorch



## **Cross-GPU operations not allowed by default**

- Exceptions: `copy_()`, `to()`, `cuda()`, and other methods with copy semantics

## **Ops on tensors across devices will cause errors**

- Can mitigate this by enabling peer-to-peer memory access

# Asynchronous Execution



**GPU operations are asynchronous by default**

**Enqueued to particular device, executed later**

**Allows execution of many more computations in parallel**

# Asynchronous Execution



**Asynchronous execution typically invisible to user**

- FIFO order of queuing
- Automatic synchronization by PyTorch between devices

# Asynchronous Execution



**Can force synchronous execution using environment variable**

**CUDA\_LAUNCH\_BLOCKING = 1**

**Useful for error handling and examining stack traces**

**Functions such as `to()`, `copy()` allow `non_blocking` argument**

# CUDA Streams



**Linear sequence of operations for execution on a single device**

**By default, each device has a default stream**

**Operations within a stream are serialized by PyTorch (order is deterministic)**

**Order of execution across streams is not deterministic**

# Device-agnostic Code



**Device-agnostic code explicitly handles GPU and CPU cases**

**Common pattern is to use argparse to read user arguments**

**Code can then be invoked with runtime flags to enable or disable CUDA**



Demo

**PyTorch and CUDA semantics**

# Summary

**Tensor operations using PyTorch tensors**

**Interoperability with NumPy**

**Understanding PyTorch support for GPUs and CUDA**

**Working with tensors on GPU-enabled devices**